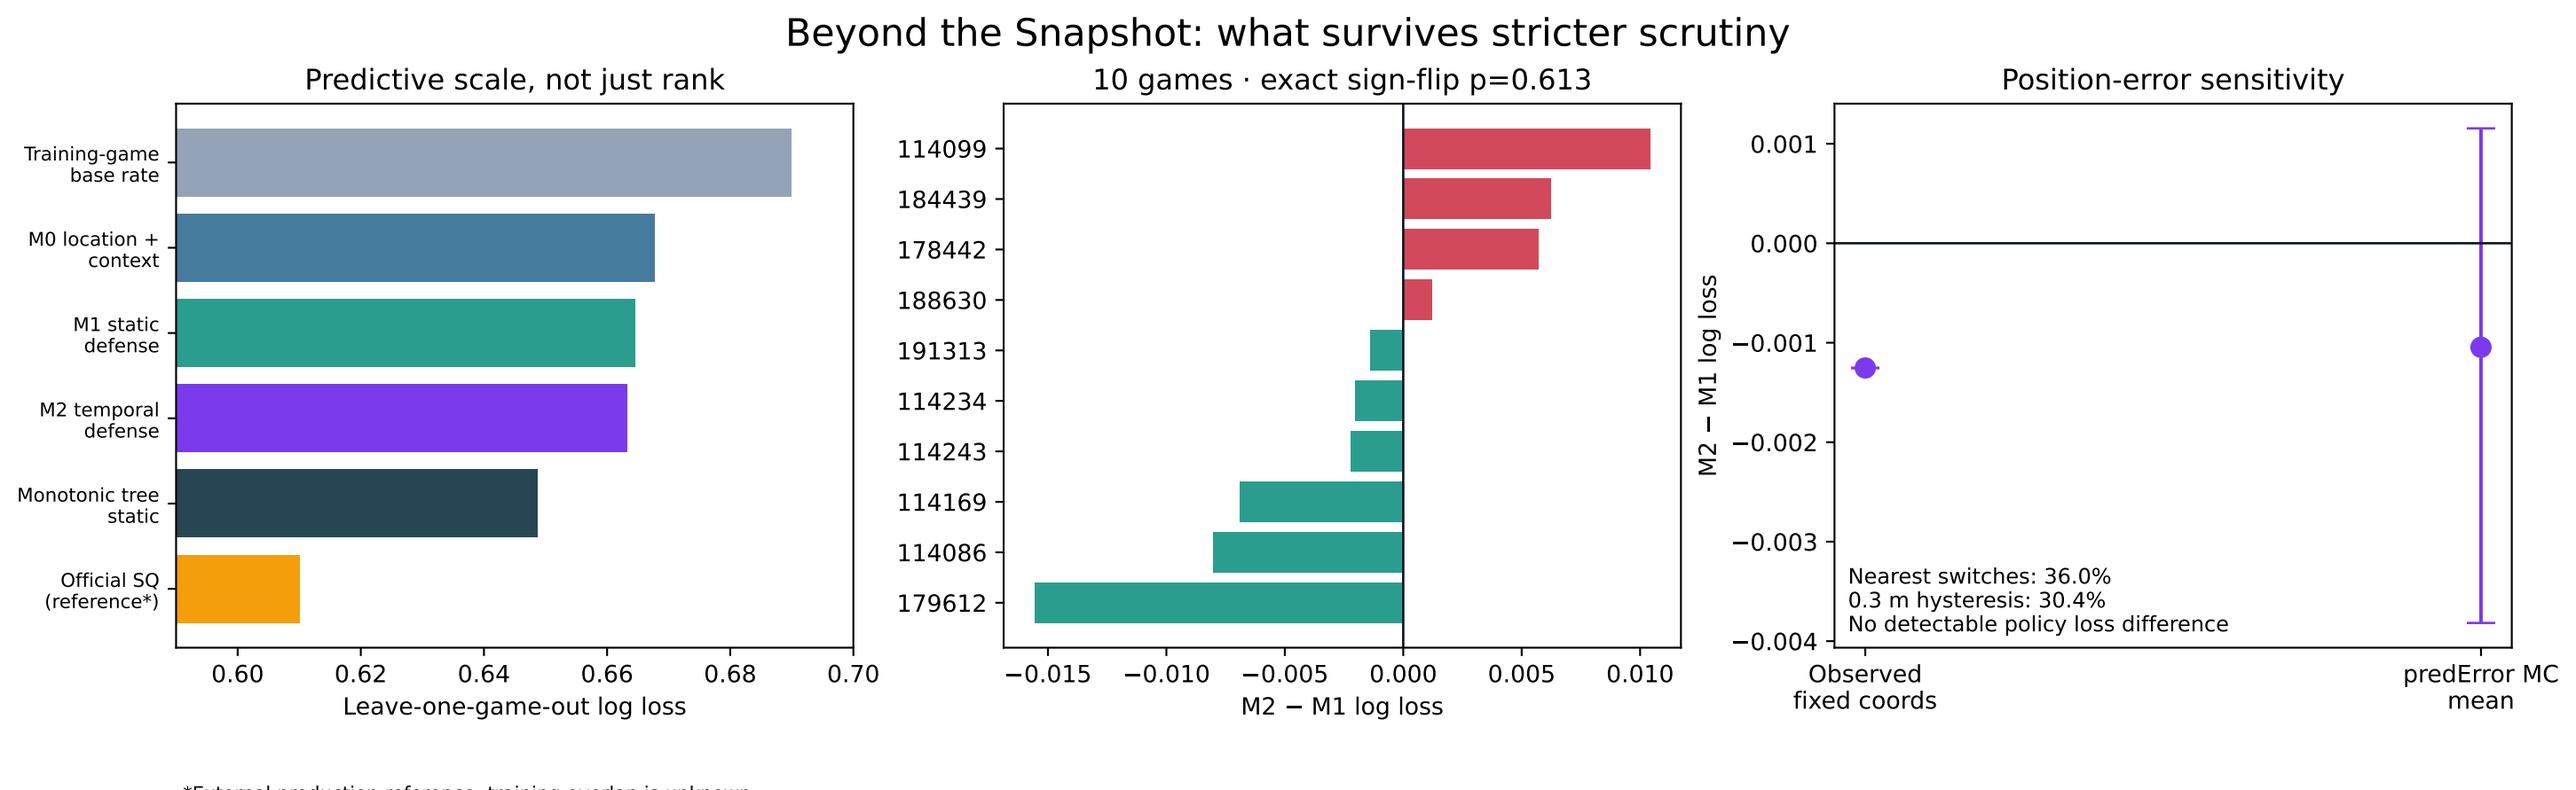
Beyond the Snapshot: what survives stricter scrutiny
Predictive scale, not just rank
Training-game base rate
M0 location + context
M1 static defense
M2 temporal defense
Monotonic tree static
Official SQ (reference*)
0.60
0.62
0.64
0.66
0.68
0.70
Leave-one-game-out log loss
10 games · exact sign-flip p=0.613
114099
184439
178442
188630
191313
114234
114243
114169
114086
179612
−0.015
−0.010
−0.005
0.000
0.005
0.010
M2 − M1 log loss
Position-error sensitivity
0.001
0.000
−0.001
−0.002
−0.003
−0.004
M2 − M1 log loss
Nearest switches: 36.0%
0.3 m hysteresis: 30.4%
No detectable policy loss difference
Observed fixed coords
predError MC mean
*External production reference; training overlap is unknown.